



SecurityEval Confusion Matrix



SecurityEval benchmark · 89 vulnerable + 89 clean Python snippets · CWE-matched

Derived Metrics at SecurityEval
Operating Point



High FPR reflects SecurityEval's synthetic clean snippets (adversarial by design).
Production adversarial corpus (Layer B, 30 repos): Confirmed Tier → 96.4% precision.