



SecurityEval Confusion Matrix



Derived Metrics at SecurityEval Operating Point



High FPR reflects SecurityEval's synthetic clean snippets (adversarial by design).
Production adversarial corpus (Layer B, 30 repos): Confirmed Tier → 96.4% precision.