

Full Evaluation Metrics — ACR-QA v5.0.0rc2 vs. Baselines

Tool / Mode	Corpus	Prec.	TPR (Recall)	F ₁	F ₃	MCC	Youden J	CVE Recall
ACR-QA Full Output	SecurityEval (89+89)	54.7%	91.0%	68.3%	0.854	0.210	0.157	8/8 (100%)
ACR-QA Confirmed Tier (P4)	Adversarial 30-repo	96.4%	37.1%	98.2%	—	—	—	8/8 (100%)
ACR-QA + LLM Augmented	RealVuln 2026	87.4%	32.4%	—	—	—	—	—
ACR-QA (Layer A core)	DVPWA/Pygoat/VulPy/DSVW	97.1%	100%†	—	—	—	—	—
Bandit (standalone)	SecurityEval / Adv.	14.0%	51.0%	21.8%	—	—	0.090	1/8 (12.5%)
Semgrep CE	SecurityEval / Adv.	36.0%	24.0%	45.7%	—	—	0.056	5/8 (62.5%)
Snyk Code	RealVuln 2026	—	17.4%	—	—	—	—	—
SonarQube CE	RealVuln 2026	—	6.5%	—	—	—	—	—

† Recall limited to annotated categories (DVPWA 50%/67% at category/file level). — = not measured on this corpus. SecurityEval: 89+89 Python snippets. Adversarial: 30-repo corpus, conservative triage.